

Comparative Study of Post-Hoc Methods vs Intrinsic Interpretability in Vision Models

Fouepe Dylan

University of Florence

June 2026

The "Black Box" Problem

- Deep learning models (CNNs, ViTs) have opaque decision-making processes.
- **Post-Hoc XAI**: Explainers applied *after* training (e.g., **SHAP** [?], **LIME** [?]).
- **Intrinsic XAI**: Models interpretable by design (e.g., **ProtoPNet** [?]).

Mathematical Foundations: XAI Methods

SHAP (Shapley Values)

$$\phi_i = \sum_S \frac{|S|!(n - |S| - 1)!}{n!} \Delta v$$

LIME (Local Surrogate)

$$\xi(x) = \arg \min_{g \in G} \mathcal{L} + \Omega(g)$$

GradCAM

$$L_{Grad-CAM}^c = ReLU \left(\sum_k \alpha_k^c A^k \right)$$

ProtoNet Similarity

$$sim = \log \left(\frac{d^2 + 1}{d^2 + \epsilon} \right)$$

Mathematical Foundations: Metrics

Infidelity

$$INFD = \mathbb{E}_I \left[\left(f(x) - f(x - I) - I^T \Phi(x) \right)^2 \right]$$

Complexity (Entropy)

$$C(\Phi) = - \sum_i p_i \ln(p_i)$$

Localization (IoU)

$$Loc = \frac{|M \cap B|}{|M \cup B|}$$


Rank Agreement (Spearman)

$$\rho = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)}$$

Experimental Setup & Metrics


Post-Hoc Metrics

XAI Metrics Comparison



Intrinsic Metrics

XAI Metrics Comparison




Datasets: *CLE4EVR* (3D objects) and *MNMath* (Handwritten equations).


Key Metrics: *Infidelity*, *Complexity*, *Localization*.

Post-Hoc Methods on CLE4EVR Dataset

ResNet50



Vision Transformer (ViT)



- **GradCAM** performs well: large 3D objects match its coarse receptive field.
- **LIME** underperforms: default SLIC superpixels severely fragment the smooth 3D shapes.

LIME Ablation Study on CLE4EVR

Hypothesis: Does matching LIME's superpixel count to the number of objects improve explanations?


- **Default** ($K = 50$): Fragments smooth 3D shapes.
- **Matched** ($K \approx 4$): Visually aligns with boundaries, but *Infidelity doubles* (from 16.0 to 34.2).

Conclusion: Linear surrogates fail to capture the non-linear reasoning of CNNs over large objects, even with perfect segmentation.

Table: LIME Segment Ablation Results


Superpixels (K)	Alignment	Infidelity ↓
Default (50)	Fragmented	16.010
Match Objects (4)	Object-Aligned	34.250
Dense (100)	Highly Fragmented	12.100

Post-Hoc Methods on MNMath Dataset




- **SHAP** heavily outperforms GradCAM on fine-grained stroke detection.
- GradCAM's low spatial resolution blurs over the thin handwritten digits.

Visual Comparison on CLE4EVR



SHAP isolates edges (sparse), while **ProtoNet Internal** highlights broad, context-aware regions.




Visual Comparison on MNMath



ProtoPNet reasoning encompasses the whole digit and its context, completely ignoring the strict "stroke" pixels favored by SHAP.


Rank Agreement: CLE4EVR Dataset

- Correlation between Spatial Locality (Intersection) and Infidelity.
- Across all architectures, post-hoc rank correlation is **near zero or negative**.



Rank Agreement: MNMATH Dataset

- The disconnect is even worse on fine-grained data (-0.24 for ProtoPNet).
- Post-hoc methods misrepresent the actual broad spatial reasoning of the network!



Key Takeaways

- Relying strictly on "Localization Accuracy" is dangerous: it assumes networks reason using human-defined object boundaries.
- Intrinsic interpretability reveals that models learn broad spatial contexts, unlike what post-hoc methods suggest [?].

Thank You!

Questions?

References I